

* BUILDER'S NOTEBOOK

Expanding the discovery layer: more sources, better classification, and a weekly generator

Allternit Blog · Builder's Notebook · July 17, 2026

BY ALLTERNIT BLOG · RESEARCH DESK

JULY 17, 2026 · 13 MIN READ

■ What shipped

We broadened the feed surface

- Added 12 publisher RSS feeds to the discovery pipeline.
- Added Bluesky and Mastodon as discovery sources.

This gives us more raw signal to work with, from traditional RSS publishers to decentralized social platforms.

We improved content classification

- Added a robotics taxonomy to categorize incoming content.

This helps route robotics-related items correctly instead of relying on keyword matching alone.

We automated the weekly update

- Added a weekly blog generator that reads the week's repository activity and produces the Builder's Notebook post.

This post was generated by that new tooling.

■ Why it matters

...

That's a bit short. Let me expand.

I need to make sure I don't use marketing superlatives. I can use "useful", "important", "helps", etc.

Let me think about the actual system. Allternit Blog seems to be a product with a discovery layer. The v1.2.11 release updated README to reflect Allternit Platform branding. This is context but predates the week.

So the blog post is about this week's commit. The commit is large - "feat(discovery): 12 pub RSS feeds, robotics taxonomy, bluesky+mastodon, weekly blog generator". It's a big feature addition.

I should write it as if we are the engineering team. "We shipped a single discovery feature this week..."

Let me write a more detailed draft:

■ Expanding the discovery layer: more sources, better classification, and a weekly generator

This week we shipped a single feature commit that expands the discovery layer in four directions at once: more RSS feeds, a new taxonomy, social sources, and a weekly blog generator. The changes touch how we collect, classify, and communicate about the platform.

■ What shipped

More RSS feeds and social sources

- Added 12 publisher RSS feeds to the discovery pipeline.
- Added Bluesky and Mastodon as discovery sources.

The discovery pipeline now pulls from a wider set of publishers plus two decentralized social networks. This increases the volume of raw content we can classify and route, while also diversifying the sources beyond traditional RSS feeds.

A robotics taxonomy

- Added a robotics taxonomy to classify incoming discovery items.

Instead of relying on keyword matching or heuristics, we now have a structured taxonomy for robotics content. This lets us tag, filter, and route robotics-related items more consistently. It also gives downstream consumers a clearer signal when they want robotics-specific content.

A weekly blog generator

- Added a generator that reads the last seven days of repository activity and produces a Builder's Notebook post.

This post is the first output of that generator. The tooling reads commit history, groups changes into themes, and drafts the post structure. We then review and publish it. The goal is to reduce the manual work of pulling the weekly log together, while keeping the engineering voice intact.

■ Why it matters

For users of the platform, the wider source surface means more content flows into the discovery system. The RSS expansion and social integrations are particularly useful for fast-moving domains like robotics, where announcements and papers often appear first on social feeds or niche publisher sites. The new taxonomy means that content gets classified more reliably, so users interested in robotics see relevant items without wading through miscategorized posts.

For the engineering team, the weekly blog generator closes a loop. We have been tracking repository activity manually for these posts; now the generator pulls the data, groups it, and produces a draft. This post is the test of that pipeline. If it works, the weekly notebook becomes a repeatable artifact rather than a writing task.

That's about 350 words. I need 600-900. I need to expand each section with more detail and perhaps some discussion of the engineering decisions or trade-offs.

But I must avoid inventing claims. So I can only elaborate on what the commit says. For example:

- 12 pub RSS feeds: I can discuss the implication of adding more feeds (increased volume, need for deduplication, etc.) but I shouldn't claim specific publishers or deduplication features unless they're in the commit.
- Robotics taxonomy: I can discuss the value of structured classification vs keyword matching.
- Bluesky+Mastodon: I can discuss the value of social APIs, rate limits, etc., but only in general terms since the commit doesn't specify implementation details.
- Weekly blog generator: I can discuss the pipeline and how it produces this post.

I can also discuss the commit as a single unit and why we chose to bundle these changes. But that might be speculative. The commit shows they were shipped

together, but I shouldn't invent the reason. I can say "the commit groups these changes" and "they share the discovery theme" without inventing a decision rationale.

Let me expand:

■ Expanding the discovery layer: more sources, better classification, and a weekly generator

This week we shipped one discovery-focused commit that touches four areas: feed ingestion, taxonomy, social sources, and the weekly notebook itself. The changes are all in the `discovery` feature area and represent a step toward a more automated and comprehensive view of what's happening across the platform.

■ What shipped

More RSS feeds and social sources

- Added 12 publisher RSS feeds to the discovery pipeline.
- Added Bluesky and Mastodon as discovery sources.

The discovery pipeline now pulls from a broader set of publishers and from two decentralized social networks. The RSS expansion increases the volume of long-form content we can track, while Bluesky and Mastodon give us a faster signal for announcements, discussions, and links that may never appear in a formal feed. Together, these sources make the ingestion layer less dependent on any single publisher or format.

A robotics taxonomy

- Added a robotics taxonomy to classify incoming discovery items.

We added a structured taxonomy for robotics content rather than relying on keyword matching or heuristics. This means incoming items can be tagged consistently, filtered by category, and routed to the right consumers. It also makes downstream features—like topic-specific feeds or search facets—easier to build, because the classification is explicit rather than inferred.

A weekly blog generator

- Added a generator that reads the last seven days of repository activity and produces a Builder's Notebook post.

This post is the first output of that generator. The tool reads the commit history, groups changes into themes, and drafts the standard post structure. We still review the output before publishing, but the generator handles the mechanical work of pulling activity, counting words, and organizing sections.

■ Why it matters

For people using the platform, the wider source surface means more content flows into discovery. The combination of RSS feeds and social sources is particularly useful for fast-moving technical domains like robotics, where papers, product updates, and discussions often appear first on social networks or specialized publisher sites. The robotics taxonomy then helps that content reach the right readers without relying on brittle keyword searches.

For the engineering team, the weekly blog generator closes a loop that has been manual until now. We have been writing these notebook posts by hand each week, which meant pulling commits, deciding what to highlight, and formatting the post. The generator turns that into a repeatable pipeline: it reads the commit data, follows the same structure, and produces a draft. This post is the first test of that pipeline. If it holds up, the Builder's Notebook becomes a low-overhead artifact of the work itself rather than a separate task.

That's about 450 words. Still need more. I can add a bit more depth to each theme, perhaps mention the single-commit nature more, and discuss the fact that v1.2.11 was the latest release but that was in April. However, the prompt says "treat as context only." I shouldn't make a big deal about it.

I could mention that the latest release remains v1.2.11 from April, so this commit is on top of that release and hasn't been tagged yet. But I should be careful not to invent release plans. I can mention it as context if relevant.

Actually, I can say: "This work lands on top of v1.2.11, which updated the README to reflect the Allternit Platform branding." That provides context but doesn't claim a new release.

Let me think about word count. Target 600-900. I need to be more expansive. Let me write a fuller version.

■ Potential additional content:

- The single commit suggests these changes are tightly coupled in the discovery pipeline. I can discuss why grouping them makes sense (ingestion, classification, output).
- Mention the impact of the blog generator on transparency: "weekly posts from actual commit data".
- Discuss the robotics taxonomy in more detail: explicit categories, consistency, maintainability.
- Discuss the social integrations: Mastodon and Bluesky are decentralized, so we need to handle different APIs, rate limits, etc. (But I shouldn't claim we handled them unless the commit says so. I can mention it as a general engineering consideration for these platforms.)
- Mention that the 12 RSS feeds are from "publishers" but not specify which ones.

Wait, I should be careful: the commit says "12 pub RSS feeds". "pub" could mean "publisher" or "publication". It's reasonable to interpret as publisher RSS feeds. I'll say "12 publisher RSS feeds".

Let me draft a fuller version:

■ Expanding the discovery layer: more sources, better classification, and a weekly generator

This week we shipped a single `discovery` feature commit that touches four parts of the pipeline at once: feed ingestion, content classification, social sources, and the weekly notebook itself. The commit is large because these pieces fit together: the more sources we ingest, the more we need structured classification; the more structured the data, the easier it is to generate a useful weekly summary.

■ What shipped

More RSS feeds and social sources

- Added 12 publisher RSS feeds to the discovery pipeline.
- Added Bluesky and Mastodon as discovery sources.

We expanded the ingestion layer to include more publisher RSS feeds and two decentralized social networks. The RSS feeds give us a steady stream of long-form articles and announcements, while Bluesky and Mastodon catch shorter, faster-moving signals. This combination means we are not relying on a single type of source or a single platform. It also means the ingestion layer needs to handle different formats and update frequencies: RSS feeds are typically polled, while social APIs have their own rate limits and content shapes.

A robotics taxonomy

- Added a robotics taxonomy to classify incoming discovery items.

We added a structured taxonomy for robotics content rather than relying on keyword matching or ad-hoc heuristics. The taxonomy gives us a consistent way to tag, filter, and route robotics-related items. That matters because "robotics" cuts across hardware, software, AI, and research; a keyword-only approach would either miss relevant items or over-include unrelated ones. With an explicit taxonomy, downstream consumers can trust the category labels and build features like topic-specific feeds or search filters without re-implementing their own classifiers.

A weekly blog generator

- Added a generator that reads the last seven days of repository activity and produces a Builder's Notebook post.

This post is the first output of that generator. The tool reads the commit log, groups changes into themes, and drafts the standard sections: what shipped, why it matters, and what's next. A human still reviews and publishes the result, but the

generator removes the manual work of pulling activity and formatting the post. The idea is to keep the engineering voice while making the notebook a byproduct of the work rather than a separate writing task.

■ Why it matters

For users of the platform, the wider source surface means more content flows into discovery. The combination of RSS feeds and social sources is especially useful for fast-moving technical domains like robotics, where papers, product updates, and discussions often appear first on social networks or specialized publisher sites. The robotics taxonomy then helps that content reach the right readers without relying on brittle keyword searches.

For the engineering team, the weekly blog generator is the other half of the change. We have been writing these notebook posts manually each week, which meant pulling commits, deciding what to highlight, and formatting the structure. The generator turns that into a repeatable pipeline: it reads the same commit data we already track, follows the same post structure, and produces a draft. This post is the first test of that pipeline. If it holds up, the Builder's Notebook becomes a low-overhead artifact of the work itself rather than a separate task.

Let me count words. Roughly: Intro: ~60 What shipped intro: ~50 More RSS feeds: ~110 Robotics taxonomy: ~110 Weekly blog generator: ~90 Why it matters: ~230

Total: ~650. Good, within 600-900.

But I need to check for marketing superlatives. I used "especially useful", "trust", "right readers". These are okay. I should avoid "revolutionary", "amazing", "best-in-class", etc. I didn't use those.

I should also check invented claims. I mentioned:

- "RSS feeds are typically polled, while social APIs have their own rate limits and content shapes." This is general knowledge, not an invented claim about our system. It's a reasonable inference. But is it too speculative? The commit doesn't mention rate limits or polling. However, it's a general engineering fact about these sources. I think it's acceptable.

- "A human still reviews and publishes the result." This is inferred from the fact that the blog post is produced by the generator and published. It could be fully automated, but the prompt says "you write this week's builder's notebook post" and it's from the repository activity. The generator produces it, but we review. Is this invented? The commit says "weekly blog generator" but doesn't say human review. I should be careful. I can say "the generator produces a draft" without claiming human review. But the fact that this post is here and coherent suggests human review or editing. Hmm. The user asks me to write the post, so presumably I am the human or the generator. But the prompt says "weekly blog generator" which is a tool. I should not claim "a human still reviews" if the commit doesn't say that. I can say "the generator produces a draft" and leave it at that.

Actually, the prompt is: "write this week's builder's notebook post (July 17, 2026) from the repository activity below." So I am acting as the writer. The commit says "weekly blog generator". So I am presumably the generator or the reviewer. The post should be honest. I can say "This post is the first output of that generator." I don't need to say human review.

Let me remove "A human still reviews and publishes the result" and instead say "The generator removes the manual work of pulling activity and formatting the post." That's more neutral.

Also, I should avoid saying "the more sources we ingest, the more we need structured classification" because that could be interpreted as the rationale for the commit. The commit groups them together but doesn't state the rationale. I can phrase it as "These pieces are connected: the ingestion changes increase the volume and variety of content, and the taxonomy helps organize it." That's a factual connection, not an invented claim.

Also, the prompt says "Only if the commit data suggests in-progress work; otherwise omit this section entirely." I should omit "What's next" because there's no in-progress work suggested.

ABOUT THIS EDITION

- Added 12 publisher RSS feeds to the discovery pipeline.

Read it on the web: allternit.com/blog

ALLTERNIT BLOG · EDITION 2026-W29

LICENSE: CC BY 4.0

GENERATED BY THE ALLTERNIT PUBLICATIONS PIPELINE